

Reading Societies Without Their Stories

Oct 2, 2026 · @Kari McKern

A multi-model test of the Complex Adaptive Model of Societies (CAMS) on ancient collapse and recovery. Working paper, neuralnations.org.

Abstract

Four structural findings about ancient collapse held across 22 blind scoring passes from three AI model families (Claude, Grok, GPT), scoring 20 society-windows on the eight CAMS nodes.

CAMS reads a society through eight functions (Helm, Shield, Lore, Stewards, Craft, Hands, Archive, Flow), each scored for Coherence, Capacity, Stress and Abstraction, without using the society's own story about itself. We scored twelve societies from Longshan China to Aboriginal Australia, plus eight extra windows around the Egyptian and Aegean collapses. Five hypotheses were registered before scoring and tested by fixed scripts.

The findings that held across all three models:

1. In both collapses, every function lost surplus (Capacity minus Stress) and shared thinking (Abstraction $\times$ Coherence) at once.
2. The Aegean recovered far more slowly than it fell; Egypt rebuilt in under three centuries.
3. In Egypt, coherence between rulers and meaning-makers fell for two centuries while surplus was still positive. The Aegean showed no such early warning.
4. Debt-relief edicts in Mesopotamia did not coincide with lower labour stress.

The models also disagreed, and the disagreement is a result in itself. They differ in scale (GPT scores about one point above Claude) and in what counts as evidence. GPT declined to score religion and records for non-literate societies, treating oral tradition as no evidence. The one hypothesis that depends on how non-state societies are judged ("no ladder") passed under Claude, landed exactly on the cut-off under Grok, and could not be tested under GPT.

1 Introduction

CAMS treats a human society as a coordination system, not a moral story. The question it asks is structural: how well does each function hold together, perform, withstand strain and think?

The model starts from a simple observation. Primate groups scaled to millions solve the problem of coordinating strangers through layers of abstraction: myth, law, writing, money. Those abstractions create an order one level up that members feel but cannot see. Its shape is set by setting and extraction strategy (river valley, coast, steppe), not by the virtue of the people in it. Like a dolphin pod or a wolf pack, a society is read by what it does, without praise or blame.

Two composite measures carry the argument:

- **Cognition** = Abstraction × Coherence: how much shared symbolic thinking a function does.
- **Energy** = Capacity – Stress: whether the function has the surplus to do it.

If CAMS is a real instrument rather than a narrative in disguise, two things must be true. Its scores should reproduce when different, independent scorers apply it. And its findings should survive a change of scorer whose background knowledge differs.

Deep time is a hard test of both. The record is thin, often written by victors or later traditions, and readily filled with modern assumptions about states and progress. This paper reports what happened when three AI model families were given the same blind rubric and asked to score the same ancient societies.

2 Method

2.1 Cases

Prompt A scored twelve society-windows at one anchor year each: Longshan (–2300), Shang Anyang (–1200), Mature Harappan (–2300), Old Kingdom Egypt (–2500), First Intermediate Egypt (–2150), Babylon under Hammurabi (–1760), Neopalatial Crete (–1600), Mycenaean LH IIIB (–1250), the Aegean after the palace destructions (–1175), pre-contact Aboriginal Australia (1700), Māori Aotearoa (1500) and the Xiongnu (–170).

Prompt B added eight windows to make time series: Egypt at –2400, –2250 and –1900; Mycenaean LH IIIA (–1350); Greece at –1000 and –750; Ur III (–2050) and late Old Babylonian (–1650). It rescored the two collapses as overlap anchors.

2.2 Rubric

Scoring used the CAMS RAW SCORER v1.2-OPT with a deep-time adapter (DT-1). Each node must pass a three-part evidence gate (an indicator of performance, an indicator of strain, a mechanism tying them to that node), or it is scored NA. Scores below 5 or above 8 need explicit functional evidence. A source posture rule treats conquerors' accounts, later traditions and colonial ethnography as reports, not as evidence of function. Scorers return CSV only, with no reasoning shown.

2.3 Independence

Every pass ran in a fresh context that saw only the prompt: no project files, no hypotheses, no other scores, and no knowledge that other passes existed. Ensembles were aggregated by a fixed script following the CAMNATIONSM5N specification, with NA handled per dimension and a mean reported only when at least three passes scored that cell.

2.4 Pre-registration

Hypotheses H1–H5 were written before any scoring. The series tests for H2, H3 and H4 were written into a script and its SHA-256 hash logged before the extra windows were scored. One later change (reporting NA cases as "not testable" rather than "against") is logged with its date.

2.5 Scorers

Model	Passes per prompt	Conditions
Claude (Opus 5.5)	15, in three independent runs of 5	Fresh subagent or conversation per pass; A and B separate
Grok 4.7	4	One conversation per prompt, prompt file only
GPT (ChatGPT Work/Codex, version unverified)	3	Fresh isolated agent per pass; A and B in the same context

An earlier, non-blind Claude pass and an earlier Grok pass that had seen the hypotheses are reported in the run log but excluded from the analysis. One upload was found to be a byte-identical copy of an earlier run and was not counted.

3 Reliability

Within each model the instrument is stable. Between models the order of societies largely holds, but the scale does not.

Comparison	Mean absolute difference per score (A / B)	Rank agreement of societies (A / B)	Mean offset
Claude run vs Claude run (5-pass ensembles)	0.14–0.23 / 0.11–0.22	0.97–0.99 / 1.00	under ± 0.15
Single Claude pass vs other Claude runs	0.26–0.29	—	—
Grok (4) vs Claude (15)	0.39–0.53 / 0.41–0.57	0.92 / 1.00	+0.2 to +0.47
GPT (3) vs Claude (15)	0.89–1.64 / 0.84–1.80	0.81 / 0.99	+0.6 to +1.2; Abstraction +1.6 to +1.8
GPT (3) vs Grok (first 3)	0.59–1.40 / 0.51–1.45	0.85 / 0.99	GPT higher, mostly on Abstraction

The spread between passes of the same model was similar for all three (standard deviation 0.12–0.44 per dimension). So the gaps between models are not noise from a single pass; they are consistent differences in how each model applies the rubric.

The ten time windows in prompt B were ranked identically or almost identically by every model (0.98–1.00). The twelve societies in prompt A, which mix states and non-state societies, were ranked less consistently (0.81–0.92 across models). The rank shifts fall on societies without visible rulers or without writing: Aboriginal Australia (GPT ranks it five places lower than Claude), the Xiongnu, Harappa, Minoan Crete and Māori Aotearoa.

Scale offsets matter less than they look for the hypotheses: most tests compare societies or windows within one model, so a model that scores everything a point higher still produces the same verdict.

4 Hypothesis results

Four of the eight tested claims held under all three models, two held in direction, and two depended on the scorer.

Hypothesis	Claude (15)	Grok (4)	GPT (3)	Status
H2a Collapse: Energy and Cognition fall together in every node	8/8 Egypt, 8/8 Aegean	8/8, 7/8	8/8, 7/8	Holds across models
H2b Recovery is slower than the fall	Aegean yes, Egypt no	Aegean yes, Egypt no	Aegean yes, Egypt yes	Aegean holds across models; Egypt mostly not
H3 Helm–Lore coherence falls before breakdown	Egypt yes, Aegean no	Egypt yes, Aegean no	Egypt yes, Aegean no	Same split across models
H4 Debt-relief edicts keep labour stress lower	fails	fails	fails	Fails across models
H1b Maritime societies: Flow above Helm	Minoan +4.1, Māori +1.1 (others –0.7)	+2.1, +1.2 (–0.6)	+2.7, +0.8 (–0.8)	Holds in direction; size varies
H1c Steppe: Shield first, Flow high	Flow 3rd	Flow 5th	Flow 3rd	Depends on the scorer
H1a River valleys: stronger Stewards and Archive	pooled +0.8 (one of three runs against)	+0.4	+0.7	Holds in direction; small and NA-sensitive
H5 No ladder: non-state societies not at the low end	holds	on the cut-off (0.798 vs 0.80)	not testable	Depends on the scorer

H4 deserves a note. Relief edicts were issued because labour was under debt strain, so higher labour stress in edict-issuing states is what one would expect, and the test as designed cannot separate cause from response. It is still reported as a failure, because that is what the pre-registered test showed.


H5 under Grok sits exactly on the line: with four passes, non-state societies had 0.798 of the state societies' mean Energy, against a threshold of 0.80 set within the project (on Cognition the ratio is 0.92). That is indeterminate, not a refutation. Under GPT it could not be tested, because half the Aboriginal nodes were NA. The verdict therefore turns on how each model treats societies without writing (section 6).

5 Collapse and recovery

The two collapses look alike on the way down and different on the way back.

Both collapses drive surplus below zero in all three models

Mean Energy (Capacity minus Stress) across eight nodes; Egypt rebounds within 250 years, the Aegean takes centuries



Ensemble means: Claude 15 passes, Grok 4, GPT 3 · runs/*/series/node_metrics.csv · 10 windows

In Egypt, mean Energy falls step by step through the late Old Kingdom, reaches about zero by -2250 (GPT keeps it slightly positive) and bottoms out in the First Intermediate Period. All three models then put the Middle Kingdom back near the level of the Old Kingdom's height. In the Aegean, the fall is abrupt: one window of strain (LH IIIB), then the deepest drop in the data. By -750, Claude and Grok still put the Aegean below its LH IIIB level; GPT puts it slightly above.

The early-warning signal in H3 is a different measure: coherence between Helm and Lore. In Egypt it fell from about 8 to between 5 and 6.5 across the two pre-collapse windows in every model, while Energy was still positive. In the Aegean it stayed flat until the breakdown. Egypt's collapse began as a loss of shared purpose at the top; the Aegean's struck a system that still looked coherent at the top.

6 What a model counts as evidence

The NA cells, where a model declined to score a node, show each model's idea of what counts as evidence, and they differ more than the scores do.

Node left NA (passes out of total)	Claude (15)	Grok (4)	GPT (3)
Harappan Helm (no visible rulers)	12	0	1
Harappan Shield	14	1	3

Node left NA (passes out of total)	Claude (15)	Grok (4)	GPT (3)
Aboriginal Australia: Helm, Shield, Lore, Archive	0	0	3 each
Māori Lore and Archive	0	0	3 each
Xiongnu Lore / Archive	5 / 7	0 / 3	3 / 3
Longshan Archive (no writing)	15	4	3
Aegean Archive after the collapse (prompt A / prompt B)	0 / 4 (otherwise scored 1)	4 / 4	2 / 2

Three patterns stand out:

- **Claude hides Harappa's rulers.** It accepts the uniform weights, bricks and seals as evidence of records and craft, but finds no mechanism linking them to a ruling centre. Grok assigns that coordination to Helm.
- **GPT does not accept oral tradition.** It leaves Lore and Archive unscored for every society without writing, and four of the eight nodes for Aboriginal Australia. Claude shows a milder version of the same caution for the Xiongnu, whose institutions are known mainly through Chinese records. Claude and Grok score Aboriginal Lore and Archive at 7–8, treating songlines and law as functioning institutions. The rubric states that oral records count, but it does not say so loudly enough for every model.
- **"Absent" and "unknown" are confused.** When a function demonstrably disappeared (literacy in the Aegean after the palaces), some passes score it 1 and others score it NA. These mean different things and the rubric does not yet separate them.

This matters beyond method. The one hypothesis about non-state societies (H5) changed verdict between models largely because of these choices. CAMS was built to read societies without borrowing their own or their rivals' stories. The models bring stories of their own, about writing, states and what counts as evidence, and the NA pattern is where those stories show.

7 Limitations

Agreement between AI models is evidence that the instrument is applied consistently, not evidence that the scores are historically true.

- **Shared corpus.** All three models learned from overlapping literatures. Where those literatures share a bias, the models will agree on it. No human specialist has yet scored any case.

- **Unequal samples.** Claude has 15 passes per prompt, Grok 4 and GPT 3. The third Claude run (R6) was supplied without a run record, and its model is inferred from its scores, not confirmed.
- **Unequal isolation.** Claude and Grok scored prompts A and B in separate contexts. GPT scored both in one context per pass, so its overlap windows and series are less independent.
- **Single anchors.** Each society is scored at one representative year. The windows stand for periods of 50 to 250 years and hide change within them.
- **Thin series.** H2 and H3 rest on two collapse series; H4 on one Mesopotamian series of three windows. Three series are too few to generalise from.
- **Choices made within the project.** The test operationalisations (for example the 0.80 threshold in H5) were written after the hypotheses and after the first, non-blind pass, though before any blind scoring. The H4 test cannot separate cause from response.
- **No reasoning shown.** The blind prompt forbids explanation to keep scorers independent, so only the first, non-blind pass has a node-by-node evidence log. Disagreements can be located but not yet explained from the scorers' own reasons.

8 Next steps

1. **Grok passes 5–6.** Two more blind Grok passes are in progress. After pass 4, Grok's H5 result sits exactly on the cut-off and its steppe Flow ranking has fallen to 5th; the last passes will show whether either moves.
2. **Rubric DT-2.** State explicitly that oral tradition is evidence for Lore and Archive. Separate "function demonstrably absent" (score low) from "no evidence" (NA). Add guidance on Abstraction for coerced but skilled labour, where models disagree most.
3. **Rescore under DT-2** with all three models, to test whether the NA differences and H5 converge once the rule is explicit.
4. **A fourth model.** Kimi, scoring blind from the prompts alone before it sees any results.
5. **More series.** Add collapse and recovery series for Mesopotamia after Ur III, the Shang-Zhou transition, the Late Bronze Age Levant and the post-Harappan Indus, to test H3 beyond Egypt and a post hoc hypothesis (H6): that recovery is faster where the setting forces re-coordination, as in a single river valley.
6. **Human scorers.** Ask specialists in two or three regions to score their cases on the same rubric, as the first test against something other than a language model.
7. **Portraits.** Write the planned "what it felt like to belong" sketch for each node of each case, now that robust scores exist to reverse-engineer them from.

Data and code

All prompts, raw scores, scripts, ensembles and the full run log are in the `backstory` folder of the `wintermute` repository (GitHub `KaliBond/wintermute`), to be published through `neuralnations.org`. Every raw file carries a SHA-256 hash.

Item	Path	SHA-256 (first 16)
Prompt A, 12 societies	<code>prompts/scorer_prompt_deeptime_v1.2-OPT.txt</code>	64009fc68f5a86b0
Prompt B, 10 windows	<code>prompts/scorer_prompt_deeptime_v1.2-OPT_windows2.txt</code>	90b683991da9128e
Series tests (pre-registered)	<code>scripts/hypothesis_check_series.py</code>	353aaeab26419658 (as registered)
Run log, runs 1–10	<code>RUNLOG.md</code>	—
Claude 15-pass reference	<code>runs/claude_grand/</code>	—
Grok 3-pass ensemble	<code>runs/grok-4.7_blind/ensemble3/</code>	—
GPT 3-pass ensemble	<code>runs/gpt_3pass/</code>	—
Full bundle	<code>Backstory_for_Kimi_2026-10-02.zip</code>	a9e08138f695cd22

The series-test script was changed once after registration (NA reported as “not testable”); the change is logged in `RUNLOG.md`. Raw scores were produced by the models from their training knowledge; the Claude and Grok passes made no tool or web calls, and the GPT run record does not state whether web access was available.

Full tables, every window and model, and node-level data: [Complete results](#)

Complete results

This tab holds every result produced so far, as of 2 October 2026: 22 blind passes per prompt across three model families, 20 society-windows, 8 nodes and 4 dimensions.


Run	Model	Passes per prompt	Used in results	Notes
1	Claude	1	No	In-session, not blind; superseded
2 and 3	Claude	5	Yes	Blind subagents; prompt A (run 2) and prompt B (run 3)
4	Grok 4.7	1	No	Had read the hypotheses before scoring
5	Claude Opus 5.5	5	Yes	Independent replication, separate session
6	Claude (inferred)	5	Yes	Second replication; model not confirmed
7, 8, 10	Grok 4.7	4	Yes	Blind, one conversation per prompt
9	GPT (Work/Codex)	3	Yes	Isolated agents; prompts A and B in one context
—	Claude	5	No	Byte-identical copy of run 5; logged as a duplicate

In the window-level tables, each model's ensemble uses prompt A's scores for the two collapse windows that both prompts scored (First Intermediate Egypt and the Aegean c.1175). A mean is reported only when at least three passes scored that cell; otherwise it is NA.

All 20 windows

All three models rank the windows alike; GPT scores the weakest far higher

Mean Node Value ($C + K - S + 0.5A$) across the eight nodes, by window, in date order



Ensemble means per window · runs/claude_grand, runs/grok-4.7_blind/ensemble4, runs/gpt_3pass · 20 windows

The models agree most at the top (the Old Kingdom, Ur III and the Middle Kingdom sit at 13–14 for all three) and least at the bottom. For the Aegean after the collapse, Claude gives –0.1, Grok 2.4 and GPT 5.8. GPT compresses the scale: the worst windows look less bad, so its collapses are shallower.

Window-level results

Mean Energy (Capacity – Stress) and mean Cognition (Abstraction × Coherence) across the eight nodes, and the number of nodes left NA in the ensemble, for each model. Date order.

Window	Year	Energy Claude	Energy Grok	Energy GPT	Cognition Claude	Cognition Grok	Cognition GPT	NA nodes C / Gk / GPT
Old Kingdom Egypt	2500 BCE	3.95	3.51	2.79	47.1	48.8	58.5	0 / 0 / 0

Window	Year	Energy Claude	Energy Grok	Energy GPT	Cognition Claude	Cognition Grok	Cognition GPT	NA nodes C / Gk / GPT
Egypt, late 5th dynasty	2400 BCE	2.69	2.38	2.25	40.0	42.7	56.0	0/0/0
Harappan cities	2300 BCE	2.11	2.23	2.77	35.6	37.1	58.7	1/0/5
Longshan	2300 BCE	0.09	0.83	0.98	22.6	27.5	36.1	1/1/2
Egypt, late 6th dynasty	2250 BCE	-0.11	0.09	0.82	27.0	28.7	46.3	0/0/0
First Intermediate Egypt	2150 BCE	-2.98	-1.94	-0.93	15.4	17.7	34.4	0/0/0
Ur III	2050 BCE	3.15	2.84	2.28	47.6	52.8	59.1	0/0/0
Middle Kingdom Egypt	1900 BCE	3.43	3.08	2.54	47.3	51.1	60.4	0/0/0
Babylon under Hammurabi	1760 BCE	2.28	1.99	1.61	42.5	42.6	54.2	0/0/0
Late Old Babylonian	1650 BCE	-0.95	-1.56	0.19	28.5	28.8	45.8	0/0/1
Neopalatial Crete	1600 BCE	1.99	2.49	2.02	37.7	40.9	53.6	1/1/2
Mycenaean LH IIIA	1350 BCE	2.23	1.58	2.12	32.6	37.0	49.5	0/0/0
Mycenaean LH IIIB	1250 BCE	0.80	1.50	1.22	33.2	40.0	51.0	0/0/0
Shang Anyang	1200 BCE	1.65	2.04	1.77	36.8	41.9	49.9	0/0/0
Aegean after the collapse	1175 BCE	-4.65	-3.14	-1.90	10.0	13.6	27.7	0/1/1
Protogeometric Greece	1000 BCE	-1.24	-0.92	0.88	13.6	18.5	39.2	1/2/4
Late Geometric Greece	750 BCE	0.74	0.75	1.45	27.6	32.3	46.3	0/0/0
Xiongnu	170 BCE	1.15	1.56	1.66	29.2	32.5	52.6	0/1/3
Māori Aotearoa	1500 CE	1.05	1.18	0.60	34.6	35.5	41.8	0/0/2

Window	Year	Energy Claude	Energy Grok	Energy GPT	Cognition Claude	Cognition Grok	Cognition GPT	NA nodes C / Gk / GPT
Aboriginal Australia	1700 CE	2.88	2.25	1.82	41.4	42.4	48.5	0 / 0 / 4

Three things the table shows that the summary does not:

- **Late Old Babylonian (1650 BCE) is a third stressed window.** Claude and Grok put it below zero (−0.95, −1.56), close to the late 6th dynasty in Egypt. It was not designed as a collapse case and is worth a series of its own.
- **GPT's Cognition runs 10–20 points above the others everywhere,** because its Abstraction scores are about 1.7 higher. Its Cognition values should only be compared within GPT.
- **Claude's NA count of 1 for Harappa** reflects the pooled rule (a mean needs at least three passes): Helm survives in only 3 of 15 passes, Shield in 1.

Hypothesis verdicts by run

Each cell is the scripted verdict from `hypothesis_check.py` (single-window tests on Prompt A) or `hypothesis_check_series.py` (series tests on Prompts A + B). Runs 2/3 are the first Claude ensemble (5 sub-agents on A; series from the run-3 overlap join). Claude pooled = all 15 Claude passes.

Test	Claude run 2/3	Claude run 5	Claude run 6	Claude pooled (15)	Grok 4.7 (4)	GPT (3)
H1a river Stewards+Archive NV gap	−0.03 against	+1.20	+0.67	+0.82	+0.40	+0.71
H1b maritime Flow–Helm (B07, B11) vs others	4.1, 0.6 vs −1.37	4.2, 0.9 vs −1.27	3.6, 1.6 vs −1.07	4.1, 1.1 vs −0.68	2.1, 1.2 vs −0.62	2.7, 0.8 vs −0.80
H1c Xiongnu: Shield rank / Flow rank	1 / 3	1 / 3	1 / 3	1 / 3	1 / 5	1 / 3
H2 fall: nodes where Energy and Cognition both fell (Egypt / Aegean)	8 / 8	8 / 8	8 / 8	8 / 8	8 / 7	8 / 7
H2 recovery slower than fall, Egypt	supports	against	against	against	against	supports

Test	Claude run 2/3	Claude run 5	Claude run 6	Claude pooled (15)	Grok 4.7 (4)	GPT (3)
H2 recovery slower than fall, Aegean	supports	supports	supports	supports	supports	supports
H3 Helm-Lore C falls before breakdown, Egypt	-3.0 supports	-3.0 supports	-2.9 supports	-3.0 supports	-2.6 supports	-1.5 supports
H3 same, Aegean	+0.1 against	not testable (Helm NA at B16)	+0.3 against	+0.2 against	+0.3 against	0.0 against
H4 level: Old Babylonian Hands Stress vs palace mean	6.2 vs 6.73 supports	6.0 vs 6.47 supports	6.2 vs 6.47 supports	6.1 vs 6.57 supports	6.8 vs 6.50 against	7.3 vs 7.43 supports (marginal)
H4 series: Mesopotamian windows below comparator	0 of 3 against	0 of 3 against	0 of 3 against	0 of 3 against	0 of 3 against	0 of 3 against
H5 non-state / state Energy (fail below 0.80)	1.85 / 1.73 = 1.07	1.55 / 1.87 = 0.83	1.83 / 1.98 = 0.92	1.69 / 1.84 = 0.92	1.66 / 2.08 = 0.80 (0.798)	1.36 / 1.88 = 0.72, not testable

Reading the table:

- **Held under every run:** H1b direction, H2 fall (all or nearly all nodes drain together), H2 slow Aegean recovery, H3 in Egypt.
- **Failed under every run:** H3 in the Aegean and H4 as a time series. The H4 design compares Mesopotamian Hands Stress with a comparator built from other palace windows, so a relief device that works only during crises cannot show up; this is a design fault more than a refutation.
- **Holds in direction, small:** H1a. Five of six columns put river valleys ahead; the gap is under one Node Value point except in run 5, and Harappan NA moves it.
- **Scorer-dependent:** H1c (Grok ranks Flow 5th), Egyptian recovery speed (split 2 supports, 4 against), H4 as a level test (Grok against), and H5. GPT's H5 ratio sits below the cut but is marked not testable because GPT left Aboriginal Helm, Shield, Lore and Archive NA, so the non-state mean rests on two societies plus half of a third.

Node Value by window, node and model

Each cell is ensemble-mean Node Value ($C + K - S + 0.5A$) as **Claude 15 / Grok 4 / GPT 3**. A dash is NA under the pooled rule (fewer than three scorers gave a value). The full table of means for all four dimensions, 480 rows, is in `results/node_means_three_models.csv`.

Window	Year	Helm	Shield	Lore	Stewards	Craft	Hands	Archive	Flow
B04 Old Kingdom Egypt	2500 BCE	16.4/15.7/15.3	11.1/11.0/12.4	16.6/15.6/15.3	14.4/13.4/13.6	18.2/16.4/16.8	11.1/11.9/11.5	14.6/15.4/14.8	12.0/13.6/13.6
B13 Egypt 5th Dyn.	2400 BCE	11.9/12.9/13.0	8.5/10.1/12.5	15.0/14.3/14.0	11.4/11.9/11.8	15.5/13.2/14.7	10.0/10.9/11.5	14.2/13.6/15.2	10.8/10.8/14.0
B01 Longshan	2300 BCE	6.1/7.7/10.0	5.1/7.2/8.1	8.0/9.8/-	6.6/7.9/9.5	12.1/12.2/13.6	6.6/8.2/8.8	-/-/-	6.6/7.4/10.0
B03 Harappan	2300 BCE	10.6/10.4/-	-/8.5/-	9.7/11.4/-	9.6/11.2/-	15.4/15.1/16.5	10.0/10.5/10.4	9.1/10.2/-	14.1/13.9/15.3
B14 Egypt 6th Dyn.	2250 BCE	4.8/7.1/9.5	5.7/6.7/9.6	10.8/10.2/13.3	6.4/5.9/9.5	9.0/9.5/13.3	6.4/7.4/8.7	10.2/9.4/13.0	7.1/7.8/9.5
B05 First Intermediate	2150 BCE	-0.1/2.0/6.6	1.7/3.2/7.0	6.6/6.0/10.2	3.4/5.4/7.7	4.0/4.3/9.4	0.7/4.4/5.6	3.8/5.9/8.5	1.4/2.9/6.3
B19 Ur III	2050 BCE	15.2/15.8/15.2	11.2/11.8/12.0	14.1/14.6/14.3	13.8/14.4/13.3	13.2/12.6/14.0	9.0/9.9/10.8	18.3/17.6/16.2	12.4/12.8/13.4
B15 Middle Kingdom	1900 BCE	15.0/14.3/15.0	14.2/13.9/15.0	15.0/14.6/14.0	11.9/12.4/12.6	15.4/14.9/15.0	9.5/11.7/11.2	16.4/15.2/15.8	11.8/13.1/14.0
B06 Old Babylonian	1760 BCE	13.7/13.5/12.7	11.6/12.1/12.0	13.0/11.7/13.3	10.8/9.8/10.8	11.6/11.1/13.7	7.0/8.3/9.2	16.4/16.0/15.8	11.7/11.0/12.0
B20 Late Old Babylonian	1650 BCE	5.6/6.1/9.4	3.4/4.1/7.3	11.0/10.0/13.0	5.2/4.1/9.3	7.4/7.2/-	3.7/3.6/6.8	12.0/10.8/14.2	5.6/4.3/9.6
B07 Minoan Crete	1600 BCE	9.2/11.0/11.5	-/-/-	11.1/12.9/-	11.6/12.5/12.6	15.6/15.4/15.5	8.0/9.4/10.4	9.0/10.6/12.4	13.3/13.1/14.2
B16 Mycenaean LH IIIA	1350 BCE	11.1/11.1/12.5	10.8/9.6/12.0	8.4/10.8/12.2	11.0/10.8/12.5	14.0/12.6/14.3	8.0/9.3/10.0	10.6/10.6/13.0	12.8/11.3/14.0
B08 Mycenaean LH IIIB	1250 BCE	9.8/11.4/11.5	9.4/11.0/11.1	8.6/10.8/12.5	10.0/11.2/11.6	11.8/12.3/14.0	6.0/8.6/9.0	10.7/12.7/13.3	9.2/10.2/11.6
B02 Shang Anyang	1200 BCE	11.6/12.5/12.2	10.0/9.8/11.9	13.1/14.3/14.1	9.4/9.9/10.5	16.0/14.6/16.5	5.5/7.4/8.8	11.8/15.3/13.7	8.2/9.1/10.2
B09 LBA collapse Aegean	1175 BCE	-2.9/0.0/4.9	0.4/2.2/4.9	3.6/4.0/7.6	-0.6/2.4/4.9	2.9/3.7/8.5	0.0/2.2/4.5	-6.2/-/-	1.8/2.3/5.4
B17 Protogeometric	1000 BCE	2.9/5.0/-	3.1/-/-	5.6/5.8/-	3.8/4.5/9.0	7.8/10.2/13.0	3.0/4.2/9.0	-/-/-	4.9/4.4/9.5
B18 Late Geometric	750 BCE	7.3/8.3/10.5	7.2/7.0/9.3	11.2/11.6/13.3	7.1/7.3/9.6	11.4/12.5/14.0	7.2/8.7/9.7	5.9/6.8/11.0	11.0/10.8/14.0
B12 Xiongnu	170 BCE	11.2/10.7/12.8	13.8/12.6/15.0	8.2/8.0/-	9.4/10.0/11.2	9.1/10.9/-	6.8/9.2/10.5	4.7/-/-	10.9/9.9/12.3
B11 Māori Aotearoa	1500 CE	7.1/8.2/9.3	7.9/9.2/9.5	13.4/12.8/-	7.8/9.0/10.2	13.2/12.0/12.6	8.2/8.1/9.2	13.0/12.0/-	8.2/9.4/10.1
B10 Aboriginal Australia	1700 CE	10.8/10.6/-	7.8/7.4/-	16.8/15.7/-	13.6/12.2/12.4	12.0/12.3/12.8	10.7/9.5/11.1	16.1/15.3/-	12.1/12.4/12.2

What the node table shows that the window means hide:

- **Hands is the lowest or second-lowest node in almost every state window under all three models.** The labour node carries the stress of the state; non-state societies (Aboriginal, Māori) do not show that gap.
- **In collapse, Lore holds up best.** Lore has the highest Node Value in B05 under all three models, and in B09 under Claude and Grok (GPT puts Craft first). The lowest nodes are Helm, Hands and Flow in B05, and Archive and Helm in B09. The meaning-making node outlasts the command and administrative ones.

- **Late Old Babylonian (B20) has the collapse signature in miniature:** Shield, Stewards, Hands and Flow drop to 3–6 for Claude and Grok, while Lore and Archive stay above 10.
- **GPT's floor is high:** its collapse windows rarely fall below 5, which compresses the contrast every hypothesis depends on.
- **Aboriginal Australia's Lore is the highest of any window for Claude and Grok, and its Archive is near the top** (oral law and song-line memory scored as working archives); GPT leaves both NA. This single rubric choice is the largest disagreement in the data set.